

# Lecture 10 Introduction to Machine Learning and Linear Regression

## Motivating Example I: Single-variable (1D) Linear Regression

### Problem

Given the *training dataset*  $(x^{(i)} \in \mathbb{R}, y^{(i)} \in \mathbb{R}), i = 1, 2, \dots, N$ , we want to find the linear function

$$y \approx f(x) = wx + b$$

that fits the relations between  $x^{(i)}$  and  $y^{(i)}$ . So that given any new  $x^{test}$  in the **test** dataset, we can make the prediction

$$y^{pred} = wx^{test} + b$$

### Training the model

- With the training dataset, define the loss function  $L(w, b)$  of parameter  $w$  and  $b$ , which is also called **mean squared error** (MSE)

$$L(w, b) = \frac{1}{N} \sum_{i=1}^N (\hat{y}^{(i)} - y^{(i)})^2 = \frac{1}{N} \sum_{i=1}^N ((wx^{(i)} + b) - y^{(i)})^2,$$

where  $\hat{y}^{(i)}$  denotes the predicted value of  $y$  at  $x^{(i)}$ , i.e.  $\hat{y}^{(i)} = wx^{(i)} + b$ .

- Then find the minimum of loss function -- note that this is the quadratic function of  $w$  and  $b$ , and we can analytically solve  $\partial_w L = \partial_b L = 0$ , and yields

$$w^* = \frac{\sum_{i=1}^N (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^N (x^{(i)} - \bar{x})^2} = \frac{\frac{1}{N} \sum_{i=1}^N (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\frac{1}{N} \sum_{i=1}^N (x^{(i)} - \bar{x})^2} = \frac{\text{Cov}(X, Y)}{\text{Var}(X)},$$
$$b^* = \bar{y} - w^* \bar{x},$$

where  $\bar{x}$  and  $\bar{y}$  are the mean of  $x$  and of  $y$ , and  $\text{Cov}(X, Y)$  denotes the estimated covariance (or called sample covariance) between  $X$  and  $Y$  (a little difference with what you learned in statistics is that we have the normalization factor  $1/N$  instead of  $1/(N-1)$  here), and  $\text{Var}(Y)$  denotes the sample variance of  $Y$  (the normalization factor is still  $1/N$ ). This is just about convention -- in statistics, they pursue for unbiased estimator.

### Evaluating the model

- MSE: The smaller MSE indicates better performance
- R-Squared: The larger  $R^2$  (closer to 1) indicates better performance. Compared with MSE, R-squared is **dimensionless**, not dependent on the units of variable.

$$R^2 = 1 - \frac{\sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2}{\sum_{i=1}^N (y^{(i)} - \bar{y})^2} = 1 - \frac{\frac{1}{N} \sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2}{\frac{1}{N} \sum_{i=1}^N (y^{(i)} - \bar{y})^2} = 1 - \frac{\text{MSE}}{\text{Var}(Y)}$$

```

In [1]: import numpy as np

class MyLinearRegression1D:
    '''
        The single-variable linear regression estimator -- writing in the style of sklearn package
    '''

    def fit(self, x, y):
        '''
            Determine the optimal parameters w, b for the input data x and y

            Parameters
            -----
            x : 1D numpy array with shape (n_samples,) from training data
            y : 1D numpy array with shape (n_samples,) from training data

            Returns
            -----
            self : returns an instance of self, with new attributes slope w (float) and intercept b (float)
        '''

        cov_mat = np.cov(x,y,bias=True) # covariance matrix, bias = True makes the factor is 1/N -- but it doesn't matter actually, since the factor will be cancelled
        self.w = cov_mat[0,1] / cov_mat[0,0] # the (0,1) element is COV(X,Y) and (0,0) element is Var(X). (1,1) is Var(Y)
        self.b = np.mean(y)-self.w * np.mean(x)

    def predict(self,x):
        '''
            Predict the output values for the input value x, based on trained parameters

            Parameters
            -----
            x : 1D numpy array from training or test data

            Returns
            -----
            returns 1D numpy array of same shape as input, the predicted y value of corresponding x
        '''

        return self.w*x+self.b

    def score(self, x, y):
        '''
            Calculate the R-squared on the dataset with input x and y

            Parameters
            -----
            x : 1D numpy array with shape (n_samples,) from training or test data
            y : 1D numpy array with shape (n_samples,) from training or test data

            Returns
            -----
            returns float, the R^2 value
        '''

        y_hat = self.predict (x) # predicted y
        mse = np.mean((y-y_hat)**2) # mean squared error
        return 1- mse / np.var(y) # return R-squared

```

```
In [2]: import pandas as pd
house = pd.read_csv('kc_house_data.csv')
house.sample(5)
```

Out[2]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
16792	6388930420	20140805T000000	582000.0	3	2.50	2380	19860	2.0	0	0
7648	7625703945	20140701T000000	345000.0	2	1.00	1080	7775	1.0	0	0
9653	1402600110	20150226T000000	392000.0	4	2.25	2360	7733	2.0	0	0
21563	9406530090	20141020T000000	337000.0	4	2.50	2470	5100	2.0	0	0
19437	5379803386	20140801T000000	289950.0	4	1.75	1500	8400	1.0	0	0

5 rows × 11 columns

```
In [3]: house.drop(['id', 'date', 'zipcode', 'lat', 'long', 'yr_built', 'yr_renovated'], axis = 1, inplace = True)
house
```

Out[3]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above
0	221900.0	3	1.00	1180	5650	1.0	0	0	3	7	1180
1	538000.0	3	2.25	2570	7242	2.0	0	0	3	7	2170
2	180000.0	2	1.00	770	10000	1.0	0	0	3	6	770
3	604000.0	4	3.00	1960	5000	1.0	0	0	5	7	1050
4	510000.0	3	2.00	1680	8080	1.0	0	0	3	8	1680
...	...	...	...	...	...	...	...	...	...	...	...
21608	360000.0	3	2.50	1530	1131	3.0	0	0	3	8	1530
21609	400000.0	4	2.50	2310	5813	2.0	0	0	3	8	2310
21610	402101.0	2	0.75	1020	1350	2.0	0	0	3	7	1020
21611	400000.0	3	2.50	1600	2388	2.0	0	0	3	8	1600
21612	325000.0	2	0.75	1020	1076	2.0	0	0	3	7	1020

21613 rows × 12 columns


```
In [4]: X = house.iloc[:,1:].to_numpy()
y = house['price'].to_numpy()
```

```
In [5]: X.shape
```

Out[5]: (21613, 13)

```
In [6]: import matplotlib.pyplot as plt  
x = X[:,2]  
plt.scatter(x,y)
```

Out[6]: <matplotlib.collections.PathCollection at 0x7fd05deda3d0>



```
In [7]: lreg = MyLinearRegression1D() # initialize the instance of one estimator
help(lreg)
```

Help on MyLinearRegression1D in module \_\_main\_\_ object:

```
class MyLinearRegression1D(builtins.object)
|   The single-variable linear regression estimator -- writing in the style of sklearn package
|
|   Methods defined here:
|
|   fit(self, x, y)
|       Determine the optimal parameters w, b for the input data x and y
|
|       Parameters
|       -----
|           x : 1D numpy array with shape (n_samples,) from training data
|           y : 1D numpy array with shape (n_samples,) from training data
|
|       Returns
|       -----
|       self : returns an instance of self, with new attributes slope w (float) and
intercept b (float)
|
|   predict(self, x)
|       Predict the output values for the input value x, based on trained parameters
|
|       Parameters
|       -----
|           x : 1D numpy array from training or test data
|
|       Returns
|       -----
|       returns 1D numpy array of same shape as input, the predicted y value of corresponding x
|
|   score(self, x, y)
|       Calculate the R-squared on the dataset with input x and y
|
|       Parameters
|       -----
|           x : 1D numpy array with shape (n_samples,) from training or test data
|           y : 1D numpy array with shape (n_samples,) from training or test data
|
|       Returns
|       -----
|       returns float, the R^2 value
|
|   -----
|   Data descriptors defined here:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)
```


```
In [8]: lreg.fit(x,y)
```

```
In [9]: lreg.score(x,y)
```

```
Out[9]: 0.49286538652201417
```

```
In [10]: fig = plt.figure(dpi = 100)
plt.scatter(x,y)
plt.plot(x,lreg.predict(x),'r')
```

Out[10]: [



```
In [11]: from sklearn import linear_model # compare with the scikit learn package
lreg_sklearn = linear_model.LinearRegression()
lreg_sklearn.fit(x.reshape(-1,1),y) #only accept 2D-array as x
```

Out[11]: LinearRegression()

```
In [12]: print(lreg.w,lreg.b)
print(lreg_sklearn.coef_, lreg_sklearn.intercept_)

280.8066899295006 -43867.60153385543
[280.80668993] -43867.60153385544
```

```
In [13]: lreg_sklearn.score(x.reshape(-1,1),y)
```

Out[13]: 0.49286538652201417

```
In [14]: help(lreg_sklern)
```

Help on LinearRegression in module sklearn.linear\_model.\_base object:

```
class LinearRegression(sklearn.base.MultiOutputMixin, sklearn.base.RegressorMixin, LinearModel)
```

```
LinearRegression(*, fit_intercept=True, normalize=False, copy_X=True, n_jobs=None)
```

Ordinary least squares Linear Regression.

LinearRegression fits a linear model with coefficients  $w = (w_1, \dots, w_p)$  to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation.

Parameters

-----

**fit\_intercept** : bool, default=True

Whether to calculate the intercept for this model. If set to False, no intercept will be used in calculations (i.e. data is expected to be centered).

**normalize** : bool, default=False

This parameter is ignored when `fit_intercept` is set to False. If True, the regressors  $X$  will be normalized before regression by subtracting the mean and dividing by the l2-norm. If you wish to standardize, please use `:class:`sklearn.preprocessing.StandardScaler`` before calling `fit` on an estimator with `normalize=False`.

**copy\_X** : bool, default=True

If True,  $X$  will be copied; else, it may be overwritten.

**n\_jobs** : int, default=None

The number of jobs to use for the computation. This will only provide speedup for `n_targets > 1` and sufficient large problems. `None` means 1 unless in a `:obj:`joblib.parallel_backend`` context. `-1` means using all processors. See `:term:`Glossary <n_jobs>`` for more details.

Attributes

-----

**coef\_** : array of shape (n\_features, ) or (n\_targets, n\_features)

Estimated coefficients for the linear regression problem. If multiple targets are passed during the fit ( $y$  2D), this is a 2D array of shape (n\_targets, n\_features), while if only one target is passed, this is a 1D array of length n\_features.

**rank\_** : int

Rank of matrix  $X$ . Only available when  $X$  is dense.

**singular\_** : array of shape (min( $X$ ,  $y$ ),)

Singular values of  $X$ . Only available when  $X$  is dense.

**intercept\_** : float or array of shape (n\_targets,)

Independent term in the linear model. Set to 0.0 if `fit_intercept = False`.

See Also

-----

`sklearn.linear_model.Ridge` : Ridge regression addresses some of the problems of Ordinary Least Squares by imposing a penalty on the size of the coefficients with l2 regularization.

`sklearn.linear_model.Lasso` : The Lasso is a linear model that estimates sparse coefficients with l1 regularization.

`sklearn.linear_model.ElasticNet` : Elastic-Net is a linear regression model trained with both l1 and l2 -norm regularization of the coefficients.

Notes

-----

From the implementation point of view, this is just plain Ordinary



Least Squares (scipy.linalg.lstsq) wrapped as a predictor object.

#### Examples

```
-----
>>> import numpy as np
>>> from sklearn.linear_model import LinearRegression
>>> X = np.array([[1, 1], [1, 2], [2, 2], [2, 3]])
>>> # y = 1 * x_0 + 2 * x_1 + 3
>>> y = np.dot(X, np.array([1, 2])) + 3
>>> reg = LinearRegression().fit(X, y)
>>> reg.score(X, y)
1.0
>>> reg.coef_
array([1., 2.])
>>> reg.intercept_
3.0000...
>>> reg.predict(np.array([[3, 5]]))
array([16.])
```

#### Method resolution order:

```
LinearRegression
sklearn.base.MultiOutputMixin
sklearn.base.RegressorMixin
LinearModel
sklearn.base.BaseEstimator
builtins.object
```

#### Methods defined here:

```
__init__(self, *, fit_intercept=True, normalize=False, copy_X=True, n_jobs=None)
    Initialize self. See help(type(self)) for accurate signature.
```

```
fit(self, X, y, sample_weight=None)
    Fit linear model.
```

#### Parameters

```
-----
X : {array-like, sparse matrix} of shape (n_samples, n_features)
    Training data

y : array-like of shape (n_samples,) or (n_samples, n_targets)
    Target values. Will be cast to X's dtype if necessary

sample_weight : array-like of shape (n_samples,), default=None
    Individual weights for each sample

.. versionadded:: 0.17
    parameter *sample_weight* support to LinearRegression.
```

#### Returns

```
-----
self : returns an instance of self.
```

---

#### Data and other attributes defined here:

```
__abstractmethods__ = frozenset()
```

---

#### Data descriptors inherited from sklearn.base.MultiOutputMixin:

```
__dict__
    dictionary for instance variables (if defined)
```

```
__weakref__
    list of weak references to the object (if defined)
```

---

#### Methods inherited from sklearn.base.RegressorMixin:

```
score(self, X, y, sample_weight=None)
```

Return the coefficient of determination  $R^2$  of the prediction.

The coefficient  $R^2$  is defined as  $(1 - u/v)$ , where  $u$  is the residual sum of squares  $((y_{\text{true}} - y_{\text{pred}}) ** 2).sum()$  and  $v$  is the total sum of squares  $((y_{\text{true}} - y_{\text{true.mean()}}) ** 2).sum()$ .

The best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of  $y$ , disregarding the input features, would get a  $R^2$  score of 0.0.

Parameters

**X** : array-like of shape (n\_samples, n\_features)

Test samples. For some estimators this may be a precomputed kernel matrix or a list of generic objects instead, shape = (n\_samples, n\_samples\_fitted), where n\_samples\_fitted is the number of samples used in the fitting for the estimator.

**y** : array-like of shape (n\_samples,) or (n\_samples, n\_outputs)  
True values for X.

**sample\_weight** : array-like of shape (n\_samples,), default=None  
Sample weights.

Returns

**score** : float

$R^2$  of self.predict(X) wrt. y.

Notes

The  $R^2$  score used when calling ``score`` on a regressor uses ``multioutput='uniform\_average'`` from version 0.23 to keep consistent with default value of :func:`~sklearn.metrics.r2\_score`. This influences the ``score`` method of all the multioutput regressors (except for :class:`~sklearn.multioutput.MultiOutputRegressor`).

---

Methods inherited from LinearModel:

**predict**(self, X)

Predict using the linear model.

Parameters

**X** : array\_like or sparse matrix, shape (n\_samples, n\_features)  
Samples.

Returns

**C** : array, shape (n\_samples,)  
Returns predicted values.

---

Methods inherited from sklearn.base.BaseEstimator:

**\_\_getstate\_\_**(self)

**\_\_repr\_\_**(self, N\_CHAR\_MAX=700)

Return repr(self).

**\_\_setstate\_\_**(self, state)

**get\_params**(self, deep=True)

Get parameters for this estimator.

Parameters

-----  
deep : bool, default=True  
    If True, will return the parameters for this estimator and  
    contained subobjects that are estimators.

Returns

-----

params : mapping of string to any  
    Parameter names mapped to their values.

set\_params(self, \*\*params)  
    Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects  
(such as pipelines). The latter have parameters of the form  
``<component>\_\_<parameter>`` so that it's possible to update each  
component of a nested object.

Parameters

-----

\*\*params : dict  
    Estimator parameters.

Returns

-----

self : object  
    Estimator instance.

# Motivating Example II: Multi-variable Linear Regression (OLS -- Ordinary Least Square)

## Problem

Given the *training dataset*  $(x^{(i)}, y^{(i)}), i = 1, 2, \dots, N$ , this time with  $y^{(i)} \in \mathbb{R}$  and  $x^{(i)} \in \mathbb{R}^p$ , we fit the multi-variable linear function

$$y \approx \mathbf{f}(x) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p = \tilde{x}\beta,$$

$$\tilde{x} = (1, x_1, \dots, x_p) \in \mathbb{R}^{1 \times (p+1)}, \beta = (\beta_0, \beta_1, \dots, \beta_p)^T \in \mathbb{R}^{(p+1) \times 1}.$$

Here  $\beta$  is called regression coefficients, and  $\beta_0$  specially referred to intercept.

Using the whole training dataset, we can write as

$$Y = \begin{pmatrix} y^{(1)} \\ y^{(2)} \\ \dots \\ y^{(N)} \end{pmatrix} \approx \begin{pmatrix} \mathbf{f}(x^{(1)}) \\ \mathbf{f}(x^{(2)}) \\ \dots \\ \mathbf{f}(x^{(N)}) \end{pmatrix} = \begin{pmatrix} \tilde{x}^{(1)}\beta \\ \tilde{x}^{(2)}\beta \\ \dots \\ \tilde{x}^{(N)}\beta \end{pmatrix} = \begin{pmatrix} \tilde{x}^{(1)} \\ \tilde{x}^{(2)} \\ \dots \\ \tilde{x}^{(N)} \end{pmatrix} \beta = \tilde{X}\beta,$$

where

$$\tilde{X} = \begin{pmatrix} 1 & x_1^{(1)} & \dots & x_p^{(1)} \\ 1 & x_1^{(2)} & \dots & x_p^{(2)} \\ \dots & \dots & \dots & \dots \\ 1 & x_1^{(N)} & \dots & x_p^{(N)} \end{pmatrix}$$

is also called the augmented data matrix.

- **Question:** To get unknown  $\beta$ , can we directly solve the linear equation  $\tilde{X}\beta = Y$ ?
- **Answer:** Most time no, because 1) typically there are more equations than variables ( $N \gg (p+1)$ ) 2) the linear model is merely the approximation to the real mapping 3) there are noises in the data points -- it's highly possible that there is NO solution at all!
- **Strategy:** Instead of solving  $\tilde{X}\beta = Y$  exactly, we want find  $\beta$  such that  $\tilde{X}\beta$  is as close as  $Y$ .

## Training the model

- With the training dataset, define the loss function  $L(\beta)$  of parameters  $\beta$ , which is also called **mean squared error (MSE)**

$$L(\beta) = \frac{1}{N} \sum_{i=1}^N (\hat{y}^{(i)} - y^{(i)})^2 = \frac{1}{N} \sum_{i=1}^N (y^{(i)} - \tilde{x}^{(i)}\beta)^2,$$

where  $\hat{y}^{(i)}$  denotes the predicted value of  $y$  at  $x^{(i)}$ , i.e.

$$\hat{y}^{(i)} = \beta_0 + \beta_1 x_1^{(i)} + \dots + \beta_p x_p^{(i)} = \tilde{x}^{(i)}\beta.$$

Now the problem becomes

$$\min_{\beta} L(\beta),$$

i.e. find the minimizer of a multi-variable ( $p+1$  dimensions) function.

- Then find the minimum of loss function -- There are two ways, either by numerical optimization (will be introduced in discussion) or by solving linear systems (introduced below), which is also called the **normal equation** approach.

To solve the critical points, we have  $\nabla L(\beta) = 0$ .

$$\frac{\partial L}{\partial \beta_0} = 2 \sum_{i=1}^N (\tilde{x}^{(i)}\beta - y^{(i)}) = 0,$$

$$\frac{\partial L}{\partial \beta_k} = 2 \sum_{i=1}^N x_k^{(i)} (\tilde{x}^{(i)}\beta - y^{(i)}) = 0, \quad k = 1, 2, \dots, p.$$

In Matrix form, it can be expressed as (left as exercise)

$$\tilde{X}^T \tilde{X} \beta = \tilde{X}^T Y,$$

also called the **normal equation** of linear regression. The optimal parameter  $\hat{\beta} = \operatorname{argmin} L(\beta)$  is also called the ordinary least square (**OLS**) estimator in statistics community.

Then the OLS estimator can be solved as

$$\hat{\beta} = (\tilde{X}^T \tilde{X})^{-1} \tilde{X}^T Y.$$

**Geometrical Interpretation** ([https://en.wikipedia.org/wiki/Ordinary\\_least\\_squares](https://en.wikipedia.org/wiki/Ordinary_least_squares))

Denote  $\tilde{X} = (\tilde{X}_0, \tilde{X}_1, \dots, \tilde{X}_p)$ , then  $\tilde{X}\beta = \sum_{k=0}^p \beta_k \tilde{X}_k$ . We require that the residual  $Y - \tilde{X}\beta$  is vertical to the plane spanned by  $\tilde{X}_k$ , which yields

$$\tilde{X}_k^T (Y - \tilde{X}\beta) = 0, \quad k = 0, 1, \dots, p$$

**Exercise:** Check that when  $p = 1$ , the solution is equivalent to the single-variable regression.

## Prediction in Test Data

Given the new observation called  $x^{(test)}$ , we have the prediction as

$$\hat{y}^{(test)} = \hat{\beta}_0 + \hat{\beta}_1 x_1^{(test)} + \dots + \hat{\beta}_p x_p^{(test)} = \tilde{x}^{(test)} \hat{\beta}.$$

## Evaluating the model

- MSE: The smaller MSE indicates better performance
- R-Squared: The larger  $R^2$  (closer to 1) indicates better performance. Compared with MSE, R-squared is **dimensionless**, not dependent on the units of variable.

$$R^2 = 1 - \frac{\sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2}{\sum_{i=1}^N (y^{(i)} - \bar{y})^2} = 1 - \frac{\frac{1}{N} \sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2}{\frac{1}{N} \sum_{i=1}^N (y^{(i)} - \bar{y})^2} = 1 - \frac{\text{MSE}}{\text{Var}(Y)}$$

The coding of multi-variable linear regression left as the homework this week. Below we will call the function in sklearn directly.

```
In [ ]: from sklearn import linear_model # compare with the scikit learn package
lreg_sklearn = linear_model.LinearRegression()
lreg_sklearn.fit(X, y)
lreg_sklearn.score(X, y)
```

```
In [ ]: lreg_sklearn.coef_
```

```
In [ ]: lreg_sklearn.intercept_
```

## Motivating Example III: Single-variable Polynomial Regression (Special Case of Multivariable Linear Regression)

### Problem

Given the *training dataset*  $(x^{(i)}, y^{(i)}), i = 1, 2, \dots, N$ , this time with  $y^{(i)} \in \mathbb{R}$  and  $x^{(i)} \in \mathbb{R}$ , we fit the single-variable polynomial function of  $p$ -th order

$$y \approx f(x) = w_0 + w_1 x + w_2 x^2 + \dots + w_p x^p$$

**Remark:** A basic conclusion in numerical analysis is that with  $N$  points, we can have a polynomial of order  $(N-1)$  that fits every point perfectly.

### Strategy

Single-variable **polynomial regression** is a special case of multi-variable **linear** regression, because we can construct a dataset of  $p$  variables by defining each row as  $(x, x^2, \dots, x^p)$  for each observation at  $x$ .

# Machine Learning: Overview of the whole picture

Possible hierarchies of machine learning concepts:

- **Problems:** Supervised Learning (Regression, Classification), Unsupervised Learning (Dimension Reduction, Clustering), Reinforcement Learning (Not covered in this course)
- **Models:**
  - (Supervised) Linear Regression, Logistic Regression, K-Nearest Neighbor (kNN) Classification/Regression, Decision Tree, Random Forest, Support Vector Machine, Ensemble Method, Neural Network...
  - (Unsupervised) K-means, Hierarchical Clustering, Principle Component Analysis, Manifold Learning (MDS, IsoMap, Diffusion Map, tSNE), Auto Encoder...
- **Algorithms:** Gradient Descent, Stochastic Gradient Descent (SGD), Back Propagation (BP), Expectation–Maximization (EM)...

For the same **problem**, there may exist multiple **models** to describe it. Given the specific **model**, there might be many different **algorithms** to solve it.

Why there is so much diversity? The following two fundamental principles of machine learning may provide theoretical insights.

**Bias-Variance Trade-off** (<https://towardsdatascience.com/understanding-the-bias-variance-tradeoff-165e6942b229>):

Simple models -- large bias, low variance. Complex models -- low bias, large variance

**No Free Lunch Theorem** (<https://analyticsindiamag.com/what-are-the-no-free-lunch-theorems-in-data-science/#:~:text=Once%20Upon%20A%20Time,that%20they%20brought%20a%20drink>):

(in plain language) There is no one model that works best for every problem. (more quantitatively) Any two models are equivalent when their performance averaged across all possible problems. --Even true for [optimization algorithms](https://en.wikipedia.org/wiki/No_free_lunch_in_search_and_optimization) ([https://en.wikipedia.org/wiki/No\\_free\\_lunch\\_in\\_search\\_and\\_optimization](https://en.wikipedia.org/wiki/No_free_lunch_in_search_and_optimization)).

# Extensions of OLS: MLE, Regularization, Ridge Regression and LASSO

*Note: The detailed mathematical derivations below are optional material. You only need to know (for quiz/exam):*

- 1) what is the relation between MLE (most likelihood estimation) and the loss function in OLS regression (ordinary least-square)
- 2) the basic concepts of Ridge regression and LASSO ;
- 3) where does the additional regularization terms in the loss function of Ridge and LASSO come from ;
- 4) which model has the best performance on training/test dataset? (or is there any theoretical guarantee?)

## Most Likelihood Estimation (MLE) and loss function in OLS

We already know what the loss function looks like in OLS. Here we first provide a mathematical explanation of this loss function from the perspective of Most Likelihood Estimation (MLE).

Recall that in linear regression, our **model assumption** is

$$y^{(i)} = \tilde{x}^{(i)}\beta + \epsilon^{(i)}, i = 1, 2, \dots, N$$

Now we further **assume** that residuals or errors  $\epsilon^{(i)}$  are as independent Gaussian random variables with identical distribution  $\mathcal{N}(0, \sigma^2)$  which has mean 0 and standard deviation  $\sigma$ .

From the density function of Gaussian distribution, the probability to observe  $\epsilon^{(i)}$  within the small interval  $[z, z + \Delta z]$  is roughly

$$\mathbb{P}(z < \epsilon^{(i)} < z + \Delta z) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{z^2}{2\sigma^2}\right) \Delta z.$$

From the data, we know indeed  $z = y^{(i)} - \tilde{x}^{(i)}\beta$ . Therefore, given  $x^{(i)}$  as fixed, the probability density (likelihood) to observe  $y^{(i)}$  is roughly

$$l(y^{(i)} | x^{(i)}, \beta) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \tilde{x}^{(i)}\beta)^2}{2\sigma^2}\right).$$

Using the *independence* assumption, the overall likelihood to observe the response data  $y^i (i = 1, 2, \dots, N)$  is

$$\mathcal{P}(y^{(i)}, 1 \leq i \leq N | \beta, x^{(i)}) = \prod_{i=1}^N l(y^{(i)} | x^{(i)}, \beta)$$

The famous **Maximum Likelihood Estimation (MLE)** theory in statistics **assumes** that we aim to find the unknown parameter  $\beta$  that maximizes the  $\mathcal{P}(\beta; x^{(i)}, y^{(i)}, 1 \leq i \leq N)$  by treating  $x^{(i)}$  and  $y^{(i)}$  as fixed numbers.

Equivalently, as the function of  $\beta$ , we can maximize

$$\ln \mathcal{P} = \sum_{i=1}^N \ln l(y^{(i)} | \beta, x^{(i)}).$$

By removing the constants, we finally arrive at the **minimization** problem of  $L^2$  loss function (whose difference with **MSE -- mean squared error** is only up to the factor  $1/N$ )

$$L(\beta) = \sum_{i=1}^N (y^{(i)} - \tilde{x}^{(i)}\beta)^2 = \|Y - \tilde{X}\beta\|_2^2.$$

## MAP (instead of MLE) Estimation in Bayesian Statistics

**Recall** the likelihood function -- we interpret it as the probability of observing the response data, given the parameter  $\beta$  as fixed, i.e. conditional probability

$$\mathcal{P}(y^{(i)}, 1 \leq i \leq N | \beta, x^{(i)}) = \prod_{i=1}^N l(y^{(i)} | x^{(i)}, \beta)$$

Now we take a bayesian approach -- assume  $\beta$  is the random variable with **prior distribution**  $\mathcal{P}(\beta)$ . Then the **posterior distribution** of  $\beta$  given the data is

$$\mathcal{P}(\beta|x^{(i)}, y^{(i)}, 1 \leq i \leq N) \propto \mathcal{P}(\beta)\mathcal{P}(y^{(i)}, 1 \leq i \leq N|\beta, x^{(i)}).$$

The **Bayesian** estimation aims to maximize the posterior distribution. It is formally termed as **Maximum A-Posteriori Estimation (MAP)**. Note that

$$\operatorname{argmax}_{\beta} \mathcal{P}(\beta|x^{(i)}, y^{(i)}, 1 \leq i \leq N) = \operatorname{argmax}_{\beta} \ln \mathcal{P}(\beta|x^{(i)}, y^{(i)}, 1 \leq i \leq N)$$

- Case 1: The prior distribution  $\mathcal{P}(\beta_i = x) \propto \exp(-x^2)$ ,  $i \geq 1$  is Gaussian-like, and different  $\beta_i$  are independent. Now the minimization problem becomes

$$\min_{\beta} ||Y - \tilde{X}\beta||_2^2 + \lambda ||\beta||_2^2.$$

here  $||\beta||_2^2 = \sum_{i=1}^p \beta_i^2$ . This is called **Ridge Regression**.

- Case 2: The prior distribution  $\mathcal{P}(\beta_i = x) \propto \exp(-|x|)$ ,  $i \geq 1$  is double-exponential like, and different  $\beta_i$  are independent. Now the minimization problem becomes

$$\min_{\beta} ||Y - \tilde{X}\beta||_2^2 + \lambda \sum_{i=1}^p |\beta_i|$$

This is called **LASSO Regression** ([https://en.wikipedia.org/wiki/Lasso\\_\(statistics\)](https://en.wikipedia.org/wiki/Lasso_(statistics))).

In general, these additional terms are called the **regularization terms**. In statistics, regularization is equivalent to Bayesian prior. Here  $\lambda$  is the adjustable parameter in algorithm -- its choice is empirical while sometimes very important for model performance (where the word "alchemy" arises in machine learning) Roughly it controls the **complexity** of the model:

- If  $\lambda \rightarrow \infty$ , we have  $\beta_i \rightarrow 0 (i \geq 1)$  and  $\beta_0 = \bar{y}$ .
- If  $\lambda \rightarrow 0$ , it will yield the same results with OLS.

Why control the complexity? Recall the bias-variance tradeoff -- sometimes reduce the complexity of model **might** help to improve performance in test dataset.

## Algorithm consideration

The optimization for ridge regression is similar to OLS -- try to derive the analytical solution your self. The optimization for LASSO is [non-trivial](https://www.cs.ubc.ca/~schmidtm/Documents/2005_Notes_Lasso.pdf) ([https://www.cs.ubc.ca/~schmidtm/Documents/2005\\_Notes\\_Lasso.pdf](https://www.cs.ubc.ca/~schmidtm/Documents/2005_Notes_Lasso.pdf)) and is the important topic in convex optimization.

## Prediction in Test Data

Now from the same training dataset, we have three  $\beta$  estimated from three different models, namely  $\hat{\beta}^{OLS}$ ,  $\hat{\beta}^{Ridge}$ ,  $\hat{\beta}^{Lasso}$  because they are the minimizers of three different loss functions.

Given the new observation called  $x^{(test)}$ , the formal expression of predictions from different methods are the same

$$\hat{y}^{(test)} = \hat{\beta}_0 + \hat{\beta}_1 x_1^{(test)} + \dots + \hat{\beta}_p x_p^{(test)} = \tilde{x}^{(test)} \hat{\beta}.$$

The **only** difference is what  $\hat{\beta}$  we use. Of course, the corresponded prediction values are also different.

## Model Performance Evaluation

- Mean Square Error (MSE) -- the lower, the better (in test data):  $\frac{1}{N} \sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2$
- R-squared (coefficient of determination,  $R^2$ ) -- the larger, the better (in test data):  $1 - \frac{\sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2}{\sum_{i=1}^N (y^{(i)} - \bar{y})^2}$

Question: What about on the training dataset?

Conclusion: **By definition**, compared with Ridge or LASSO regression, OLS **will be sure** to have the smallest MSE (hence largest  $R^2$ ) on **training dataset**. Think why!



## Example: Diabetes Dataset

We use the [scikit-learn package](https://scikit-learn.org/stable/index.html) (<https://scikit-learn.org/stable/index.html>) to load the data and run regression. More tutorials about linear models can be [found here](https://scikit-learn.org/stable/modules/linear_model.html) ([https://scikit-learn.org/stable/modules/linear\\_model.html](https://scikit-learn.org/stable/modules/linear_model.html)).

Data from [this paper](https://web.stanford.edu/~hastie/Papers/LARS/LeastAngle_2002.pdf) ([https://web.stanford.edu/~hastie/Papers/LARS/LeastAngle\\_2002.pdf](https://web.stanford.edu/~hastie/Papers/LARS/LeastAngle_2002.pdf)) by Professor [Robert Tibshirani et al](https://statweb.stanford.edu/~tibs/index.html) (<https://statweb.stanford.edu/~tibs/index.html>).

```
In [ ]: from sklearn import datasets
X,y= datasets.load_diabetes(return_X_y = True)
```

```
In [ ]: help(datasets.load_diabetes)
```

Generate the training and test dataset by random splitting

```
In [ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
```

```
In [ ]: print(X_train.shape)
print(y_test.shape)
```

```
In [ ]: help(train_test_split)
```

Ordinary Least Square (OLS) Linear Regression

```
In [ ]: from sklearn import linear_model
reg_ols = linear_model.LinearRegression()
reg_ols.fit(X_train,y_train) # train the parameters in training dataset
```

```
In [ ]: dir(reg_ols)
```

```
In [ ]: reg_ols.coef_
```

```
In [ ]: y_pred_ols = reg_ols.predict(X_test) # generate predictions in test dataset
y_pred_ols
```

```
In [ ]: from sklearn.metrics import mean_squared_error
mse_ols = mean_squared_error(y_test, y_pred_ols)
R2_ols = reg_ols.score(X_test,y_test) # the R-squared value -- how good is the fitting in test dataset?
print(mse_ols,R2_ols)
```

```
In [ ]: reg_ridge = linear_model.Ridge(alpha=.02) # alpha is proportional to the lambda above
-- only up to the constant
reg_ridge.fit(X_train,y_train)
print(reg_ridge.coef_)

y_pred_ridge = reg_ridge.predict(X_test)
mse_ridge = mean_squared_error(y_test, y_pred_ridge)
R2_ridge = reg_ridge.score(X_test,y_test)
print(mse_ridge,R2_ridge)
```

```
In [ ]: reg_lasso = linear_model.Lasso(alpha=0.1) # alpha is proportional to the lambda above
-- only up to the constant
reg_lasso.fit(X_train,y_train)
print(reg_lasso.coef_)

y_pred_lasso = reg_lasso.predict(X_test)
mse_lasso = mean_squared_error(y_test, y_pred_lasso)
R2_lasso = reg_lasso.score(X_test,y_test)
print(mse_lasso,R2_lasso)

In [ ]: print(reg_ols.score(X_train,y_train)) # note that we calculate score on TRAINING data set
print(reg_ridge.score(X_train,y_train))
print(reg_lasso.score(X_train,y_train))
```

By definition, OLS has the smallest MSE (largest R-squared) on **training dataset**. What about on the test dataset?

```
In [ ]: import numpy as np
train_errors = list()
test_errors = list()
alphas = np.logspace(-5, -1, 20)
for alpha in alphas:
    reg_lasso.set_params(alpha=alpha) # change the parameter of reg_lasso
    reg_lasso.fit(X_train, y_train)
    train_errors.append(reg_lasso.score(X_train, y_train))
    test_errors.append(reg_lasso.score(X_test, y_test))

In [ ]: import matplotlib.pyplot as plt
fig = plt.figure(dpi=100)
plt.semilogx(alphas,train_errors,label = 'train R2')
plt.semilogx(alphas,test_errors,label = 'test R2')
plt.xlabel('alpha')
plt.legend()
```

Therefore, a good model in training dataset does not mean it's a good model in the final test dataset. Then how can we use the best model (i.e. model with best regularization parameters)?

## Cross Validation ([https://scikit-learn.org/stable/modules/cross\\_validation.html](https://scikit-learn.org/stable/modules/cross_validation.html))

What if we don't know the true labels in test, but the performance in test is so important to us so that we really want to select a model with greater confidence with training dataset?

As discussed previously, we can use training dataset to make 10 "quizzes" (each "quiz" is called a validation dataset), and let the three models to compete based on the 10 "competitions". This is called 10-fold cross-validation.

For the more detailed discussion and distinguishment between training, validation and test datasets, you can refer to this [wikipedia link \(https://en.wikipedia.org/wiki/Training\\_validation\\_and\\_test\\_sets#Test\\_dataset\)](https://en.wikipedia.org/wiki/Training_validation_and_test_sets#Test_dataset).

```
In [ ]: from sklearn.model_selection import cross_val_score
scores_lasso = cross_val_score(reg_lasso, X_train, y_train, cv=10) # cross-validation function in sklearn
scores_ridge = cross_val_score(reg_ridge, X_train, y_train, cv=10)
scores_ols = cross_val_score(reg_ols, X_train, y_train, cv=10)

In [ ]: print(scores_lasso)
print(scores_ridge)
print(scores_ols)

In [ ]: help(cross_val_score)
```

```
In [ ]: import pandas as pd
scores_all = pd.DataFrame({"lasso": scores_lasso, "ols": scores_ols, "ridge": scores_ri
dge})
scores_all
```

Besides mean and standard deviation, we can also use the [boxplot \(https://towardsdatascience.com/understanding-boxplots-5e2df7bcbd51\)](https://towardsdatascience.com/understanding-boxplots-5e2df7bcbd51) to visualize the results.

```
In [ ]: import seaborn as sns
sns.set_theme()
fig, ax = plt.subplots(dpi=100)
sns.boxplot(data = scores_all)
```

```
In [ ]: scores_all.describe()
```

Of course, the final judgement is still in the test dataset.

```
In [ ]: reg_lasso.score(X_test,y_test)
```

```
In [ ]: reg_ridge.score(X_test,y_test)
```

```
In [ ]: reg_ols.score(X_test,y_test)
```

Exercise: Use cross-validation to select the `alpha` parameter in LASSO

```
In [ ]: # your code here
```

## Reference Reading Suggestions

- ISLR: Chapter 2,3,6
- ESL: Chapter 1,2,3
- PML: Chapter 1,2,3,4,7,11
- DL: Chapter 5